Exercise Sheet 2 - Solutions*

Omar D. Domingues
omar (dot) darwiche-domingues (at) inria.fr

Last update: January 19, 2019

1 Exercise 1 (Finite-Horizon Problems and Dynamic Programming)

1.1 Notation

A stochastic and non-stationary policy $\pi = (\pi_1, \pi_2, \ldots, \pi_{T-1})$ is such that $\pi_t(a|s) = P(A_t = a|S_t = s)$ represents the probability of choosing the action a in state s at time t.

If the policy π_t is deterministic, that is, $\pi_t(\bar{a}_s|s) = 1$ for some action $\bar{a}_s$, we write π_t as a function $\pi_t : \mathcal{S} \to \mathcal{A}$ from the state space $\mathcal{S}$ to the action space $\mathcal{A}$ such that $\pi_t(s) = \bar{a}_s$.

1.2 Bellman Equation

In this exercise, we consider a finite-horizon problem. The agent must choose actions in order to maximize the sum of the rewards up to a fixed time T. The value function at time t is defined as:

$$V_t^\pi(s) = E_\pi \left[\sum_{n=t}^{T-1} r_t(S_n, A_n) + r_T(S_T) \Big| S_t = s \right] \tag{1}$$

and represents how much reward is collected by following the policy $\pi = (\pi_1, \pi_2, \ldots, \pi_{T-1})$ starting from time t. Thus, the agent seeks to find a policy π that maximizes V_1^π. Note that the reward function at the final time T does not depend on the action: this is because the interactions between the agent and the environment ends at time T, and the agent's action does not matter once the time is up.

Imagine that π is given and we wish to compute $V_1^\pi(s)$. A very naive computation is to do the following sum:

$$V_1^\pi(s) = \sum_{a_1, a_2, \ldots, a_{T-1}} \sum_{s_1, s_2, \ldots, s_T} \left(\sum_{n=1}^{T-1} r_t(s_n, a_n) + r_T(s_T) \right) P(s_1, a_1, s_2, a_2, s_3, a_3, \ldots s_{T-1}, a_{T-1}, s_T | S_1 = s_1) \tag{2}$$

Let $|\mathcal{S}|$ and $|\mathcal{A}|$ be the number of states and the number of actions, respectively. Each term in the sum above requires $\mathcal{O}(T)$ multiplications (to compute the joint probability from the conditional ones) and we have a sum over $\mathcal{O}(T|\mathcal{S}|^T|\mathcal{A}|^T)$ terms. This gives a total of $\mathcal{O}(T^2|\mathcal{S}|^T|\mathcal{A}|^T)$ operations, which is huge!

We can find a smarter way to compute it by rewriting $V_t^\pi(s)$ as follows:

*Only for the exercises solved during the course.

$$
\begin{aligned}
V_t^\pi(s) &= E_\pi\left[\sum_{n=t}^{T-1} r_t(S_n, A_n) + r_T(S_T)\Big|S_t = s\right] \\
&= E_\pi\left[r_t(S_t, A_t) + \sum_{n=t+1}^{T-1} r_t(S_n, A_n) + r_T(S_T)\Big|S_t = s\right] \\
&= E_\pi\left[r_t(S_t, A_t)\Big|S_t = s\right] + E_\pi\left[E_\pi\left[\sum_{n=t+1}^{T-1} r_t(S_n, A_n) + r_T(S_T)\Big|S_{t+1}, S_t = s\right]\Big|S_t = s\right] \\
&= E_\pi\left[r_t(S_t, A_t)\Big|S_t = s\right] + E_\pi\left[V_{t+1}^\pi(S_{t+1})\Big|S_t = s\right]
\end{aligned} \tag{3}
$$

Finally, we obtain:

$$
V_t^\pi(s) = \sum_a \pi_t(a|s)\left[r_t(s,a) + \sum_{s'} p(s'|s,a)V_{t+1}^\pi(s')\right] \tag{4}
$$

and, by definition, we have:

$$
V_T(s) = E_\pi[r_T(S_T)|S_T = s] = r_T(s) \text{ for all } s \in \mathcal{S} \tag{5}
$$

Equation 4 is called *Bellman Equation* for evaluating the policy π, which is a dynamic programming approach.

Now, we can find V_1^π by the following procedure:

1. Initialization: set $t = T$, $V_T(s) = r_T(s)$;
2. Compute V_{t-1} using equation 4;
3. Set $t \leftarrow t - 1$;
4. If $t = 1$ stop, otherwise go to step 2.

At each iteration t, equation 4 requires $\mathcal{O}(1)$ multiplications and a sum over $\mathcal{O}(|\mathcal{S}||\mathcal{A}|)$ terms to compute V_t for one state. This gives $\mathcal{O}(|\mathcal{S}|^2|\mathcal{A}|)$ operations per iteration. Since we have T iterations, we have a total of $\mathcal{O}(T|\mathcal{S}|^2|\mathcal{A}|)$ operations, which is much less than the naive approach! This is the interest of using dynamic programming.

1.3 Bellman Optimatily Equation - Intuition

In this exercise we want to prove that the optimal value function $V_t^*(s) := \max_\pi V_t^\pi(s)$ satisfies the equations:

$$
\begin{aligned}
V_t^*(s) &= \max_{a\in\mathcal{A}}\left[r_t(s,a) + \sum_{s'\in\mathcal{S}} p(s'|s,a)V_{t+1}^*(s')\right] \text{ for } 1 \le t \le T-1 \\
V_T^*(s) &= r_T(s)
\end{aligned} \tag{6}
$$

First, remember that our policy is non stationary and is written as $\pi = (\pi_1, \pi_2, \ldots, \pi_{T-1})$. Note that $V_t^\pi(s)$ depends only on the actions taken starting from time t. Let's define $\pi_{t:T-1} := (\pi_t, \pi_2, \ldots, \pi_{T-1})$. With a slight abuse of notation, we can write $V_t^\pi = V_t^{\pi_{t:T-1}}$.

Now, let's see the intuition behind equation 6. We have, by using equation 4:

$$
\begin{aligned}
V_t^*(s) &= \max_{\pi} V_t^{\pi}(s) \\
&= \max_{\pi} \sum_a \pi_t(a|s) \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{\pi}(s') \right] \\
&= \max_{(\pi_t, \pi_{t+1:T-1})} \sum_a \pi_t(a|s) \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{\pi}(s') \right] \\
&= \max_{\pi_t} \sum_a \pi_t(a|s) \left[r_t(s,a) + \max_{\pi_{t+1:T-1}} \sum_{s'} p(s'|s,a) V_{t+1}^{\pi}(s') \right] \\
&= \max_{\pi_t} \sum_a \pi_t(a|s) \left[r_t(s,a) + \max_{\pi_{t+1:T-1}} \sum_{s'} p(s'|s,a) V_{t+1}^{\pi_{t+1:T-1}}(s') \right] \\
&= \max_{a} \left[r_t(s,a) + \max_{\pi_{t+1:T-1}} \sum_{s'} p(s'|s,a) V_{t+1}^{\pi_{t+1:T-1}}(s') \right]
\end{aligned}
\tag{7}
$$

where the second line comes from equation 4 and the last line comes from the fact that $\sum_a \pi_t(a|s) f(a) \leq \sum_a \pi_t(a|s) \max_a f(a) = \max_a f(a)$ for any function f and equality can be achieved by choosing $\pi_t(a|s)$ that puts probability 1 for one action $\bar{a}$ such that $\bar{a} \in \text{argmax}_a f(a)$.

Observe that:

$$
V_{t+1}^*(s') = \max_{\pi_{t+1:T-1}} V_{t+1}^{\pi_{t+1:T-1}}(s') \tag{8}
$$

Thus, if we can invert the max and the sum operator in the last line of equation 7, we prove our desired result, which is equation 6. But can we do this formally?

1.4 Bellman Optimality Equation - Formal Proof

Let $\bar{V}_t$ be such that:

$$
\begin{aligned}
&\bar{V}_t(s) = \max_{a \in \mathcal{A}} \left[r_t(s,a) + \sum_{s' \in \mathcal{S}} p(s'|s,a) \bar{V}_{t+1}(s') \right] \text{ for } 1 \leq t \leq T-1 \\
&\bar{V}_T(s) = r_T(s)
\end{aligned}
\tag{9}
$$

We will prove that $\bar{V}_t(s) = V_t^*(s)$ for all $s \in \mathcal{S}$ and for all t in two steps:

1. Step 1: prove that $V_t^*(s) \leq \bar{V}_t(s)$
2. Step 2: prove that $\bar{V}_t(s) \leq V_t^*(s)$;

1.4.1 Step 1

Let's proceed by induction on t:

- For $t = T$, we have $\bar{V}_T(s) = V_T^*(s) = r_T(s)$;
- Assume that $\bar{V}_n^*(s) \leq \bar{V}_n(s)$ for all $n \in \{t+1, \ldots, T\}$ and for all s.

We have, for an arbitrary policy π and all s:

$$
\begin{aligned}
V_t^{\pi}(s) &= \sum_a \pi_t(a|s) \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{\pi}(s') \right] \\
&\leq \max_a \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{\pi}(s') \right] \\
&\leq \max_a \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{*}(s') \right] \\
&\leq \max_a \left[r_t(s,a) + \sum_{s'} p(s'|s,a) \bar{V}_{t+1}(s') \right] \\
&= \bar{V}_t(s)
\end{aligned}
\tag{10}
$$

Since π is arbitrary, we have $\forall \pi,\ V_t^{\pi}(s) \leq \bar{V}_t(s) \implies \max_{\pi} V_t^{\pi}(s) \leq \bar{V}_t(s)$ which is what we wanted to prove.

1.4.2 Step 2

Define a deterministic policy $\pi' = (\pi'_1, \pi'_2, \ldots, \pi'_{T-1})$ such that:

$$
\begin{aligned}
\pi'_{T-1}(s) &\in \operatorname*{argmax}_a \left[r_{T-1}(s,a) + \sum_{s'} p(s'|s,a) r_T(s') \right] \\
\pi'_t(s) &\in \operatorname*{argmax}_a \left[r_t(s,a) + \sum_{s'} p(s'|s,a) V_{t+1}^{\pi'}(s') \right]
\end{aligned}
\tag{11}
$$

It can be proven by induction that $V_t^{\pi'}(s) = \bar{V}_t(s)$ for all s. Finally:

$$
\bar{V}_t(s) = V_t^{\pi'}(s) \leq \max_{\pi} V_t^{\pi}(s) = V_t^{*}(s) \text{ for all } s \tag{12}
$$